

Untitled

November 6, 2022

```
[1]: import pandas
import seaborn
```

```
[2]: df = pandas.read_csv('Auto.csv')
print('First 5 rows:')
print(df.iloc[:5])
print('Dimensions:')
print(df.shape)
```

First 5 rows:

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	\
0	18.0	8	307.0	130	3504	12.0	70.0	
1	15.0	8	350.0	165	3693	11.5	70.0	
2	18.0	8	318.0	150	3436	11.0	70.0	
3	16.0	8	304.0	150	3433	12.0	70.0	
4	17.0	8	302.0	140	3449	NaN	70.0	

	origin	name
0	1	chevrolet chevelle malibu
1	1	buick skylark 320
2	1	plymouth satellite
3	1	amc rebel sst
4	1	ford torino

Dimensions:

(392, 9)

```
[3]: print('Describe MPG:')
print(df['mpg'].describe())
print('Describe weight:')
print(df['weight'].describe())
print('Describe year:')
print(df['year'].describe())
```

Describe MPG:

count	392.000000
mean	23.445918
std	7.805007
min	9.000000
25%	17.000000

```

50%      22.750000
75%      29.000000
max       46.600000
Name: mpg, dtype: float64
Describe weight:
count      392.000000
mean      2977.584184
std       849.402560
min       1613.000000
25%      2225.250000
50%      2803.500000
75%      3614.750000
max       5140.000000
Name: weight, dtype: float64
Describe year:
count      390.000000
mean       76.010256
std        3.668093
min        70.000000
25%        73.000000
50%        76.000000
75%        79.000000
max        82.000000
Name: year, dtype: float64

```

MPG Range: 37.6, average: 23.445918 Weight range: 3527, average: 2977.584184 Year range: 12, average: 76.010256

```
[4]: print(df.dtypes)
```

```

mpg          float64
cylinders     int64
displacement  float64
horsepower    int64
weight        int64
acceleration  float64
year          float64
origin        int64
name          object
dtype: object

```

```
[5]: df['cylinders'] = df['cylinders'].astype('category').cat.codes
```

```
[6]: df['origin'] = df['origin'].astype('category')
```

```
[7]: print('Cylinders dtype: ' + str(df['cylinders'].dtype))
      print('Origin dtype: ' + str(df['origin'].dtype))
```

```
Cylinders dtype: int8
```

Origin dtype: category

```
[8]: df = df.dropna()
      print('New dimensions:')
      print(df.shape)
```

New dimensions:
(389, 9)

```
[9]: avg_mpg = df['mpg'].describe()['mean']
      df['mpg_high'] = df['mpg'].apply(lambda x: 1 if x > avg_mpg else 0).
      ↪astype('category')
```

```
[10]: if 'mpg' in df:
        del df['mpg']
      if 'name' in df:
        del df['name']
```

```
[11]: print('First 5 rows:')
      print(df.iloc[:5])
```


First 5 rows:

	cylinders	displacement	horsepower	weight	acceleration	year	origin	\
0	4	307.0	130	3504	12.0	70.0	1	
1	4	350.0	165	3693	11.5	70.0	1	
2	4	318.0	150	3436	11.0	70.0	1	
3	4	304.0	150	3433	12.0	70.0	1	
6	4	454.0	220	4354	9.0	70.0	1	

	mpg_high
0	0
1	0
2	0
3	0
6	0

```
[12]: seaborn.catplot(df['mpg_high'])
```


```
[12]: <seaborn.axisgrid.FacetGrid at 0x1eee13395e0>
```



Comment: mpg_high has both zero and one values and nothing in between (obviously, since we defined these categories ourselves)

```
[13]: seaborn.relplot(df, x='horsepower', y='weight', hue='mpg_high')
```


```
[13]: <seaborn.axisgrid.FacetGrid at 0x1eee12c78b0>
```



Comment: High MPG vehicles seem to be more clustered towards the low end of the weight/horsepower spectrum (as expected), while there is a much greater variety of low MPG vehicles.

```
[14]: seaborn.boxplot(df, x='mpg_high', y='weight')
```

```
[14]: <AxesSubplot: xlabel='mpg_high', ylabel='weight'>
```



Comment: The box plot confirms the difference in variety between low and high MPG vehicles that was noted above.

```
[15]: import sklearn
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(df, df['mpg_high'],
    ↪ test_size=0.2, random_state=1234)
del X_train['mpg_high']
print('Train dimensions:')
print(X_train.shape)
del X_test['mpg_high']
print('Test dimensions:')
print(X_test.shape)
```

```
Train dimensions:
(311, 7)
Test dimensions:
(78, 7)
```

```
[16]: from sklearn import preprocessing
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
```

```

from sklearn.pipeline import make_pipeline
from sklearn.metrics import classification_report
pipe = make_pipeline(StandardScaler(), LogisticRegression(solver='lbfgs')) #  

    ↳Without scaling, lbfgs doesn't converge
pipe.fit(X_train, y_train)
y_pred = pipe.predict(X_test)
print(classification_report(y_test, y_pred))

```

	precision	recall	f1-score	support
0	1.00	0.82	0.90	50
1	0.76	1.00	0.86	28
accuracy			0.88	78
macro avg	0.88	0.91	0.88	78
weighted avg	0.91	0.88	0.89	78

```

[17]: from sklearn import tree
pipe = make_pipeline(StandardScaler(), tree.
    ↳DecisionTreeClassifier(random_state=1234))
pipe.fit(X_train, y_train)
y_pred = pipe.predict(X_test)
print(classification_report(y_test, y_pred))

```

	precision	recall	f1-score	support
0	0.96	0.92	0.94	50
1	0.87	0.93	0.90	28
accuracy			0.92	78
macro avg	0.91	0.92	0.92	78
weighted avg	0.93	0.92	0.92	78

```

[18]: from sklearn.neural_network import MLPClassifier

pipe = make_pipeline(StandardScaler(), MLPClassifier(solver='lbfgs',
    ↳hidden_layer_sizes=(7, 3), random_state=1234))
pipe.fit(X_train, y_train)
y_pred = pipe.predict(X_test)
print('Network 1:')
print(classification_report(y_test, y_pred))

pipe = make_pipeline(StandardScaler(), MLPClassifier(solver='lbfgs',
    ↳hidden_layer_sizes=(9, 15), random_state=1234))
pipe.fit(X_train, y_train)
y_pred = pipe.predict(X_test)

```

```
print('Network 2:')
print(classification_report(y_test, y_pred))
```

Network 1:

	precision	recall	f1-score	support
0	0.92	0.92	0.92	50
1	0.86	0.86	0.86	28
accuracy			0.90	78
macro avg	0.89	0.89	0.89	78
weighted avg	0.90	0.90	0.90	78

Network 2:

	precision	recall	f1-score	support
0	0.98	0.92	0.95	50
1	0.87	0.96	0.92	28
accuracy			0.94	78
macro avg	0.92	0.94	0.93	78
weighted avg	0.94	0.94	0.94	78

Performance was somewhat dependent on the values of `hidden_layer_sizes` but seemed to hover in the range 0.88-0.92 in many cases. Increasing the number of hidden layers would intuitively make the model better, but after a certain point it begins to overfit and performance drops. (9, 15) were the best settings I found in a reasonable amount of time.

0.1 Analysis

The (9, 15) neural network algorithm performed best, while the decision tree algorithm was almost identical in performance. For all of the accuracy, precision, and recall metrics, the (9, 15) neural network performs best, followed by the decision tree, followed by the (7, 3) neural network, and worst of all is the logistic regression algorithm. The best neural network most likely outperformed the others as neural networks are extremely powerful and can learn all sorts of functions; nevertheless, “in real life” the decision tree might be better for such a small dataset, as neural networks are prone to overfitting and are less efficient to train (though that was not noticeable here).

Python and Scikit-Learn are much more enjoyable to use than R, as Python feels more like a “real” programming language and follows many conventions similar to other languages. At the same time, Python and libraries such as Pandas add a huge amount of “syntactic sugar” that allows you to write very expressive, shorthand code if you know what you’re doing.